

# Dataset 1

In [134]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge, RidgeCV, Lasso
from sklearn.preprocessing import StandardScaler
```

In [135]:

```
data=pd.read_csv(r"C:\Users\chila\Downloads\Advertising.csv")
data
```

Out[135]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...	...	...	...	...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

In [136]:

```
data.head()
```

Out[136]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

In [137]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0    TV          200 non-null    float64
 1   Radio       200 non-null    float64
 2  Newspaper   200 non-null    float64
 3   Sales      200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

In [138]:

```
data.describe
```

Out[138]:

```
<bound method NDFrame.describe of
0    230.1    37.8    69.2    22.1
1     44.5    39.3    45.1    10.4
2     17.2    45.9    69.3    12.0
3    151.5    41.3    58.5    16.5
4    180.8    10.8    58.4    17.9
..     ...     ...     ...     ...
195   38.2     3.7    13.8     7.6
196   94.2     4.9     8.1    14.0
197  177.0     9.3     6.4    14.8
198  283.6    42.0    66.2    25.5
199  232.1     8.6     8.7    18.4

[200 rows x 4 columns]>
```

In [139]:

```
plt.figure(figsize = (10, 10))  
sns.heatmap(data.corr(), annot = True)
```


Out[139]:

<Axes: >



In [140]:

```
data.drop(columns = ["Radio", "Newspaper"], inplace = True)
#pairplot
sns.pairplot(data)
data.Sales = np.log(data.Sales)
```



In [141]:

```
features = data.columns[0:2]
target = data.columns[-1]
#X and y values
x = data[features].values
y = data[target].values
#split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=17)
print("The dimension of X_train is {}".format(x_train.shape))
print("The dimension of X_test is {}".format(x_test.shape))
#Scale features
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

The dimension of X\_train is (140, 2)

The dimension of X\_test is (60, 2)

In [142]:

```
#Model
lr = LinearRegression()
#Fit model
lr.fit(x_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(x_train, y_train)
test_score_lr = lr.score(x_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 1.0

The test score for lr model is 1.0

In [143]:

```
#Ridge Regression Model
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
#train and test scorefor ridge regression
train_score_ridge = ridgeReg.score(x_train, y_train)
test_score_ridge = ridgeReg.score(x_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```


Ridge Model:

The train score for ridge model is 0.9902871391941609

The test score for ridge model is 0.984426628514122

In [144]:

```
plt.figure(figsize = (10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='green')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



In [145]:

```
#Lasso regression model
print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(x_train,y_train)
train_score_ls =lasso.score(x_train,y_train)
test_score_ls =lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.0


The test score for ls model is -0.0042092253233847465

In [146]:

```
pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[146]:

<Axes: >



In [147]:

```
#Using the Linear CV model
from sklearn.linear_model import LassoCV
#Lasso Cross validation
lasso_cv = LassoCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10], random_state=0).fit(x_train, y_train)
#score
print(lasso_cv.score(x_train, y_train))
print(lasso_cv.score(x_test, y_test))
```

0.9999999343798134

0.9999999152638072

## ELASTIC NET REGRESSION

In [151]:

```
from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
```

```
[0.00417976 0.          ]
2.026383919311004
```

In [152]:

```
y_pred_elastic=regr.predict(x_train)
```

In [153]:

```
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

```
Mean Squared Error on test set 0.5538818050142158
```

In [154]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge, RidgeCV, Lasso
from sklearn.preprocessing import StandardScaler
```



In [155]:

```
data=pd.read_csv(r"C:\Users\chila\Downloads\fiat500_VehicleSelection_Dataset.csv")
data
```

Out[155]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	
0	1	lounge	51	882	25000	1	44.907242	8.611
1	2	pop	51	1186	32500	1	45.666359	12.241
2	3	sport	74	4658	142228	1	45.503300	11.417
3	4	lounge	51	2739	160000	1	40.633171	17.634
4	5	pop	73	3074	106880	1	41.903221	12.495
...	...	...	...	...	...	...	...	...
1533	1534	sport	51	3712	115280	1	45.069679	7.704
1534	1535	lounge	74	3835	112000	1	45.845692	8.666
1535	1536	pop	51	2223	60457	1	45.481541	9.413
1536	1537	lounge	51	2557	80750	1	45.000702	7.682
1537	1538	pop	51	1766	54276	1	40.323410	17.568

1538 rows × 9 columns



In [156]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ID              1538 non-null   int64
1   model           1538 non-null   object
2   engine_power    1538 non-null   int64
3   age_in_days     1538 non-null   int64
4   km              1538 non-null   int64
5   previous_owners 1538 non-null   int64
6   lat             1538 non-null   float64
7   lon             1538 non-null   float64
8   price           1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [157]:

```
data=data[['engine_power','price']]
data.columns=['Eng','pri']
```

In [158]:

```
data.head()
```

Out[158]:

	Eng	pri
<b>0</b>	51	8900
<b>1</b>	51	8800
<b>2</b>	74	4200
<b>3</b>	51	6000
<b>4</b>	73	5700

In [159]:

```
data.tail()
```

Out[159]:


	Eng	pri
<b>1533</b>	51	5200
<b>1534</b>	74	4600
<b>1535</b>	51	7500
<b>1536</b>	51	5990
<b>1537</b>	51	7900

In [160]:

```
sns.lmplot(x='Eng',y='pri',data=data,order=2,ci=None)
```

Out[160]:

<seaborn.axisgrid.FacetGrid at 0x260b1488490>



In [161]:

```
data.describe()
```

Out[161]:

	Eng	pri
count	1538.000000	1538.000000
mean	51.904421	8576.003901
std	3.988023	1939.958641
min	51.000000	2500.000000
25%	51.000000	7122.500000
50%	51.000000	9000.000000
75%	51.000000	10000.000000
max	77.000000	11100.000000

In [162]:

```
data.fillna(method='ffill')
```

Out[162]:

	Eng	pri
0	51	8900
1	51	8800
2	74	4200
3	51	6000
4	73	5700
...	...	...
1533	51	5200
1534	74	4600
1535	51	7500
1536	51	5990
1537	51	7900

1538 rows × 2 columns

In [163]:

```
x=np.array(data['Eng']).reshape(-1,1)
y=np.array(data['pri']).reshape(-1,1)
```

In [164]:

```
data.dropna(inplace=True)
```

C:\Users\chila\AppData\Local\Temp\ipykernel\_15008\1368182302.py:1: Setting  
WithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
data.dropna(inplace=True)
```


In [165]:

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25)
#splitting the dataset into training and testing dataset
regr=LinearRegression()
regr.fit(x_train,y_train)
print(regr.score(x_test,y_test))
```

0.07262451631405753

In [166]:

```
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.scatter(x_test,y_test,color='k')
plt.show()
```




In [167]:

```
df500=data[:][:500]
sns.lmplot(x='Eng',y='pri',data=df500,order=1,ci=None)
```

Out[167]:

<seaborn.axisgrid.FacetGrid at 0x260b1886590>



## Dataset2

In [168]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn.linear_model import Ridge
from sklearn.preprocessing import StandardScaler
```

In [169]:

```
df=pd.read_csv(r"C:\Users\chila\Downloads\bottle.csv.zip")  
df
```

C:\Users\chila\AppData\Local\Temp\ipykernel\_15008\2541466974.py:1: DtypeWarning: Columns (47,73) have mixed types. Specify dtype option on import or set low\_memory=False.

```
df=pd.read_csv(r"C:\Users\chila\Downloads\bottle.csv.zip")
```

Out[169]:

Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta
0	1	1	054.0 056.0	19-4903CR-HY-060-0930-05400560-0000A-3	0	10.500	33.4400	NaN 25.64900
1	1	2	054.0 056.0	19-4903CR-HY-060-0930-05400560-0008A-3	8	10.460	33.4400	NaN 25.65600

In [170]:

```
df.head()
```

Out[170]:

Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat
3	1	4	054.0 056.0	19-4903CR-HY-060-0930-05400560-0010A-7	19	10.450	33.4200	NaN 25.64300	
0	1	1	054.0 056.0	19-4903CR-HY-060-0930-05400560-0000A-3	0	10.50	33.440	NaN 25.649	NaN
4	1	5	054.0 056.0	19-4903CR-HY-060-0930-05400560-0000A-3	20	10.450	33.4210	NaN 25.64300	
1	...	...	054.0 056.0	19-4903CR-HY-060-0930-05400560-0008A-3	8	10.46	33.440	NaN 25.656	NaN
864858	34404	864859	093.4 026.4	20-1611SR-MX-310-02239-00040264-0000A-7	0	18.744	33.4083	5.805 23.87055	
2	1	3	054.0 056.0	19-4903CR-HY-060-0930-05400560-0010A-7	10	10.46	33.437	NaN 25.654	NaN
864859	34404	864860	093.4 026.4	20-1611SR-MX-310-02239-00040264-0000A-7	2	18.744	33.4083	5.805 23.87072	
3	1	4	054.0 056.0	19-4903CR-HY-060-0930-05400560-0019A-3	19	10.45	33.420	NaN 25.643	NaN
864860	34404	864861	093.4 026.4	20-1611SR-MX-310-02239-00040264-0000A-7	5	18.692	33.4150	5.796 23.88911	
4	1	5	054.0 056.0	19-4903CR-HY-060-0930-05400560-0020A-7	20	10.45	33.421	NaN 25.643	NaN
864861	34404	864862	093.4 026.4	20-1611SR-MX-310-02239-00040264-0000A-7	10	18.161	33.4062	5.816 24.01426	

5 rows x 74 columns





In [171]:

```
Cst_Cnt  Btl_Cnt  Sta_ID  Depth_ID  Depthm  T_degC  Salnty  O2ml_L  STheta
df.info()
```

				20-					
				1611SR-					
864862	34404	864863	093.4	MX-310-	15	17.533	33.3880	5.774	24.15297
			026.4	2239-					
				09340264-					
				0015A-3					

864863 rows × 74 columns

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 864863 entries, 0 to 864862
Data columns (total 74 columns):
```

#	Column	Non-Null Count	Dtype
0	Cst_Cnt	864863 non-null	int64
1	Btl_Cnt	864863 non-null	int64
2	Sta_ID	864863 non-null	object
3	Depth_ID	864863 non-null	object
4	Depthm	864863 non-null	int64
5	T_degC	853900 non-null	float64
6	Salnty	817509 non-null	float64
7	O2ml_L	696201 non-null	float64
8	STheta	812174 non-null	float64
9	O2Sat	661274 non-null	float64
10	Oxy_μmol/Kg	661268 non-null	float64
11	BtlNum	118667 non-null	float64
12	RecInd	864863 non-null	int64
13	T_prec	853900 non-null	float64
14	T_qual	23127 non-null	float64
15	S_prec	817509 non-null	float64
16	S_qual	74914 non-null	float64
17	P_qual	673755 non-null	float64
18	O_qual	184676 non-null	float64
19	SThtaq	65823 non-null	float64
20	O2Satq	217797 non-null	float64
21	ChlorA	225272 non-null	float64
22	Chlqua	639166 non-null	float64
23	Phaeop	225271 non-null	float64
24	Phaqua	639170 non-null	float64
25	PO4uM	413317 non-null	float64
26	PO4q	451786 non-null	float64
27	SiO3uM	354091 non-null	float64
28	SiO3qu	510866 non-null	float64
29	NO2uM	337576 non-null	float64
30	NO2q	529474 non-null	float64
31	NO3uM	337403 non-null	float64
32	NO3q	529933 non-null	float64
33	NH3uM	64962 non-null	float64
34	NH3q	808299 non-null	float64
35	C14As1	14432 non-null	float64
36	C14A1p	12760 non-null	float64
37	C14A1q	848605 non-null	float64
38	C14As2	14414 non-null	float64
39	C14A2p	12742 non-null	float64
40	C14A2q	848623 non-null	float64
41	DarkAs	22649 non-null	float64
42	DarkAp	20457 non-null	float64
43	DarkAq	840440 non-null	float64
44	MeanAs	22650 non-null	float64
45	MeanAp	20457 non-null	float64
46	MeanAq	840439 non-null	float64
47	IncTim	14437 non-null	object
48	LightP	18651 non-null	float64
49	R_Depth	864863 non-null	float64
50	R_TEMP	853900 non-null	float64
51	R_POTEMP	818816 non-null	float64
52	R_SALINITY	817509 non-null	float64
53	R_SIGMA	812007 non-null	float64
54	R_SVA	812092 non-null	float64
55	R_DYNHT	818206 non-null	float64

```
56 R_O2                696201 non-null float64
57 R_O2Sat             666448 non-null float64
58 R_SIO3              354099 non-null float64
59 R_PO4               413325 non-null float64
60 R_NO3               337411 non-null float64
61 R_NO2               337584 non-null float64
62 R_NH4               64982 non-null float64
63 R_CHLA              225276 non-null float64
64 R_PHAE0             225275 non-null float64
65 R_PRES              864863 non-null int64
66 R_SAMP              122006 non-null float64
67 DIC1                1999 non-null float64
68 DIC2                224 non-null float64
69 TA1                 2084 non-null float64
70 TA2                 234 non-null float64
71 pH2                 10 non-null float64
72 pH1                 84 non-null float64
73 DIC Quality Comment 55 non-null object
```

dtypes: float64(65), int64(5), object(4)  
memory usage: 488.3+ MB

In [172]:

```
df.describe()
```

Out[172]:

	Cst_Cnt	Btl_Cnt	Depthm	T_degC	Salnty	O
count	864863.000000	864863.000000	864863.000000	853900.000000	817509.000000	696201.0
mean	17138.790958	432432.000000	226.831951	10.799677	33.840350	3.3
std	10240.949817	249664.587269	316.050259	4.243825	0.461843	2.0
min	1.000000	1.000000	0.000000	1.440000	28.431000	-0.0
25%	8269.000000	216216.500000	46.000000	7.680000	33.488000	1.3
50%	16848.000000	432432.000000	125.000000	10.060000	33.863000	3.4
75%	26557.000000	648647.500000	300.000000	13.880000	34.196900	5.5
max	34404.000000	864863.000000	5351.000000	31.140000	37.034000	11.1

8 rows × 70 columns



In [173]:

```
df.isna().any()
```

Out[173]:

```
Cst_Cnt          False
Btl_Cnt          False
Sta_ID           False
Depth_ID         False
Depthm           False
...
TA1              True
TA2              True
pH2              True
pH1              True
DIC Quality Comment  True
Length: 74, dtype: bool
```

In [174]:

```
df.isnull().sum()
```

Out[174]:

```
Cst_Cnt          0
Btl_Cnt          0
Sta_ID           0
Depth_ID         0
Depthm           0
...
TA1              862779
TA2              864629
pH2              864853
pH1              864779
DIC Quality Comment  864808
Length: 74, dtype: int64
```

In [175]:

```
df=df[['Salnty', 'T_degC']]
df.columns=['Sal', 'Temp']
```

In [176]:

```
df.head(20)
```

Out[176]:


	Sal	Temp
0	33.440	10.50
1	33.440	10.46
2	33.437	10.46
3	33.420	10.45
4	33.421	10.45
5	33.431	10.45
6	33.440	10.45
7	33.424	10.24
8	33.420	10.06
9	33.494	9.86
10	33.510	9.83
11	33.580	9.67
12	33.640	9.50
13	33.689	9.32
14	33.847	8.76
15	33.860	8.71
16	33.876	8.53
17	NaN	8.45
18	33.926	8.26
19	33.980	7.96

In [177]:

```
sns.lmplot(x='Sal',y='Temp',data=df,order=2,ci=None)
```

Out[177]:

<seaborn.axisgrid.FacetGrid at 0x260b18867a0>



In [178]:

```
df.fillna (method='ffill')
```

Out[178]:

	Sal	Temp
0	33.4400	10.500
1	33.4400	10.460
2	33.4370	10.460
3	33.4200	10.450
4	33.4210	10.450
...	...	...
864858	33.4083	18.744
864859	33.4083	18.744
864860	33.4150	18.692
864861	33.4062	18.161
864862	33.3880	17.533

864863 rows × 2 columns

In [179]:

```
df.fillna(value=0,inplace=True)
```

C:\Users\chila\AppData\Local\Temp\ipykernel\_15008\709118144.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
df.fillna(value=0,inplace=True)
```

In [180]:

```
x=np.array(df['Sal']).reshape(-1,1)
y=np.array(df['Temp']).reshape(-1,1)
```

In [181]:

```
df.dropna (inplace=True)
```

C:\Users\chila\AppData\Local\Temp\ipykernel\_15008\1939022369.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
df.dropna (inplace=True)
```

In [182]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
```


In [183]:

```
regr=LinearRegression()  
regr.fit(x_train,y_train)  
print(regr.score (x_test,y_test))
```

0.014223284614636067

In [184]:

```
y_pred=regr.predict(x_test)  
plt.scatter(x_test,y_test,color='b')  
plt.plot(x_test, y_pred, color='k')  
plt.show()
```






In [185]:

```
df500=df[:] [:500]  
sns.lmplot(x='Sal', y= 'Temp', data=df500,order=1, ci=None)
```

Out[185]:


<seaborn.axisgrid.FacetGrid at 0x260b1347a90>



In [186]:

```
df500.fillna (method='ffill', inplace=True)
x=np.array(df500['Sal']).reshape(-1,1)
y=np.array(df500['Temp']).reshape(-1,1)
df500.dropna (inplace=True)
x_train,x_test,y_train, y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression: ", regr.score (x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test, y_pred, color='k')
plt.show()
```

Regression: 0.053364290956062765



In [187]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model = LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.053364290956062765

In [188]:

```
#conclusion: Linear regression is not fit for the model
```

In [189]:

```
features = df.columns[0:2]
target = df.columns[-1]
#X and y values
x = df[features].values
y = df[target].values
#split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=17)
print("The dimension of X_train is {}".format(x_train.shape))
print("The dimension of X_test is {}".format(x_test.shape))
#Scale features
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

The dimension of X\_train is (605404, 2)  
The dimension of X\_test is (259459, 2)

## Elastic Net

In [190]:

```
from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
```

[0. 0.94934511]  
0.5401219631068042

In [191]:

```
y_pred_elastic=regr.predict(x_train)
```

In [192]:

```
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

Mean Squared Error on test set 114.40984808659205

## Dataset 3

In [210]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge, RidgeCV, Lasso
from sklearn.preprocessing import StandardScaler
```

In [211]:

```
d=pd.read_csv(r"C:\Users\chila\Downloads\fiat500_VehicleSelection_Dataset.csv")
d
```

Out[211]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	
0	1	lounge	51	882	25000	1	44.907242	8.611
1	2	pop	51	1186	32500	1	45.666359	12.241
2	3	sport	74	4658	142228	1	45.503300	11.417
3	4	lounge	51	2739	160000	1	40.633171	17.634
4	5	pop	73	3074	106880	1	41.903221	12.495
...	...	...	...	...	...	...	...	...
1533	1534	sport	51	3712	115280	1	45.069679	7.704
1534	1535	lounge	74	3835	112000	1	45.845692	8.666
1535	1536	pop	51	2223	60457	1	45.481541	9.413
1536	1537	lounge	51	2557	80750	1	45.000702	7.682
1537	1538	pop	51	1766	54276	1	40.323410	17.568

1538 rows × 9 columns



In [212]:

```
d.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   ID                    1538 non-null   int64
1   model                 1538 non-null   object
2   engine_power          1538 non-null   int64
3   age_in_days           1538 non-null   int64
4   km                    1538 non-null   int64
5   previous_owners       1538 non-null   int64
6   lat                   1538 non-null   float64
7   lon                   1538 non-null   float64
8   price                 1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [214]:

```
d.head()
```

Out[214]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	153800
1	2	pop	51	1186	32500	1	45.666359	12.241890	153800
2	3	sport	74	4658	142228	1	45.503300	11.417840	153800
3	4	lounge	51	2739	160000	1	40.633171	17.634609	153800
4	5	pop	73	3074	106880	1	41.903221	12.495650	153800

In [215]:

```
d.tail()
```

Out[215]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
1533	1534	sport	51	3712	115280	1	45.069679	7.7041	153800
1534	1535	lounge	74	3835	112000	1	45.845692	8.6661	153800
1535	1536	pop	51	2223	60457	1	45.481541	9.4131	153800
1536	1537	lounge	51	2557	80750	1	45.000702	7.6821	153800
1537	1538	pop	51	1766	54276	1	40.323410	17.5681	153800

In [216]:


```
d.drop(columns=["model"],inplace=True)
```

In [217]:

```
plt.figure(figsize=(10,10))  
sns.heatmap(d.corr(),annot = True)
```

Out[217]:

<Axes: >



In [219]:

```
features = d.columns[0:2]
target = d.columns[-1]
#X and y values
x = d[features].values
y = d[target].values
#split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=17)
print("The dimension of X_train is {}".format(x_train.shape))
print("The dimension of X_test is {}".format(x_test.shape))
#Scale features
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

The dimension of X\_train is (1076, 2)  
The dimension of X\_test is (462, 2)

In [220]:

```
#Model
lr = LinearRegression()
#Fit model
lr.fit(x_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(x_train, y_train)
test_score_lr = lr.score(x_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.07448634159905865  
The test score for lr model is 0.07913288661070894

In [221]:


```
#Ridge Regression Model
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(x_train, y_train)
#train and test score for ridge regression
train_score_ridge = ridgeReg.score(x_train, y_train)
test_score_ridge = ridgeReg.score(x_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.07448028989896427  
The test score for ridge model is 0.07885996726883049

In [222]:

```
plt.figure(figsize = (10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='green')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```






In [223]:

```
pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[223]:

<Axes: >



In [224]:

```
from sklearn.linear_model import LassoCV
#Lasso Cross validation
lasso_cv = LassoCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10], random_state=0).fit(x_train, y_train)
#score
print(lasso_cv.score(x_train, y_train))
print(lasso_cv.score(x_test, y_test))
```

0.07448634159905387

0.07913288806451946

## Elastic net

In [225]:

```
from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
```

```
[ 8.46751882e-02 -1.30405006e+02]
15279.442735227916
```

In [226]:

```
y_pred_elastic=regr.predict(x_train)
```

In [227]:

```
mse=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mse)
```

```
Mean Squared Error on test set 48390222.80186546
```

In [ ]: